

强化学习基础算法实验报告

姓名: _____ 张韫译萱_____ 学号: _____08023214_____

一、 实验题目




图1 状态转移概率模型

图 1: 状态转移模型

1 实验任务

1. 实现值迭代与策略迭代(模型 P 已知),对三种(非终止状态)奖赏值 $R(S) = \{0.01, -0.01, -0.04\}$, 找出最优策略。
2. 实现 Exploratory-MC 与 Q-Learning (模型 P 未知), 非终止状态的奖赏值为 $R(s) = -0.04$, 并实验对比它们的性能。

2 实验要求

1. 选择 C++ 或 Python 实现。
2. 代码以文本形式粘贴在附录相应位置, 注释准确, 能成功运行。

二、 实验结果

1 实验参数设置

本实验中自定义的参数如下: 折扣因子 $\gamma = 0.9$, 该值能够平衡即时奖励与长期收益; 收敛阈值 $\theta = 10^{-4}$, 作为值迭代和策略评估的终止条件; Q-Learning 学习率 $\alpha = 0.1$; 探索率 $\epsilon = 0.1$; MC 和 Q-Learning 均训练 10000 轮。

2 值迭代与策略迭代运行结果

对三种不同的每步奖赏值 $R(s)$ 进行实验，得到的最优策略如下。箭头表示该状态下的最优动作，背景颜色深浅代表状态价值。




图 2: $R(s) = 0.01$ 时的最优策略与价值函数





图 3: $R(s) = -0.01$ 时的最优策略与价值函数

图 4: $R(s) = -0.04$ 时的最优策略与价值函数

3 Exploratory-MC 与 Q-Learning 运行结果

在 $R(s) = -0.04$ 的情况下，分别使用蒙特卡洛（Monte Carlo ES）和 Q-Learning 算法进行学习。




图 5: Monte Carlo ES 学习到的策略




图 6: Q-Learning 学习到的策略




图 7: Monte Carlo ES 与 Q-Learning 学习曲线对比

三、实验分析

1 值迭代与策略迭代分析

实验结果表明，步奖励 $R(s)$ 的设置对最优策略具有显著影响。

当 $R(s) = 0.01$ 时，非终止状态具有正奖励，智能体倾向于延长轨迹长度以累积更多奖励。策略图中可观察到部分状态的动作指向远离终止状态的方向，这是因为在折扣因子 $\gamma = 0.9$ 的作用下，持续获取正奖励的累积收益可能超过直接到达 +1 终止状态的收益。

当 $R(s) = -0.01$ 时，每步存在轻微惩罚，智能体开始倾向于寻找通往 +1 终止状态的路径。但

由于惩罚较小，智能体会优先选择风险较低的路径，避免靠近-1 终止状态。例如，(2,1) 位置的智能体选择向上移动而非向右，以降低误入-1 状态的概率。

当 $R(s) = -0.04$ 时，步惩罚增大，智能体更倾向于选择最短路径到达 +1 终止状态。此时路径长度的代价超过了规避风险的收益，因此策略变得更加激进。这一设置产生的策略与经典 Grid World 问题的最优解一致。

值迭代和策略迭代虽然实现方式不同，但收敛到相同的最优策略。值迭代直接对状态值函数进行 Bellman 最优更新，而策略迭代采用策略评估与策略改进交替进行的方式。

2 免模型算法分析

MC 和 Q-Learning 均为无模型方法，无需已知环境的状态转移概率，通过与环境交互进行学习。

Q-Learning 采用时序差分更新，每执行一步即可更新 Q 值。其更新公式为 $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ ，其中 $\max_{a'} Q(s', a')$ 项使得算法具有 off-policy 特性，即学习的是最优策略而非当前执行的策略。实验结果显示 Q-Learning 收敛较快且曲线相对平稳。

Monte Carlo ES 需要等待完整 episode 结束后才能更新 Q 值，采用首次访问的方式计算状态-动作对的回报均值。”探索起点”机制通过随机选择初始状态和动作来保证所有状态-动作对都能被充分访问。由于回报 G 的方差较大且更新频率较低，MC 的学习曲线波动明显大于 Q-Learning。

两种算法最终均能收敛到接近最优的策略。Q-Learning 在本问题中表现更优，主要原因是 Grid World 的 episode 较短，单步更新的效率优势明显。

3 收敛性与时间分析




图 8: 四种算法的收敛过程与时间对比

从收敛图可以分析四种算法的特性。

值迭代的 delta 值（状态值最大变化量）呈指数下降趋势，约 20 余次迭代即可收敛至 10^{-4} 以下。这是由于 Bellman 更新算子是压缩映射，收敛速度由 γ 决定。总耗时约为数毫秒级别。

策略迭代仅需 3-4 轮即可找到最优策略。每轮包含一次完整的策略评估（需要多次迭代）和一次策略改进。由于策略空间远小于值空间，且策略改进保证单调性，策略迭代在本问题上收敛更快。

Q-Learning 的 Q 值变化在前 1000 个 episode 内快速下降并趋于稳定。单步更新机制使得信息利用效率较高。后期曲线不会完全归零，这是由于 ϵ -greedy 策略持续引入探索噪声。总耗时约为数秒。

Monte Carlo 的收敛速度明显慢于 Q-Learning，且波动更大。这是 MC 方法的固有特点：必须等待 episode 结束才能更新，且回报 G 的方差较大。从时间曲线可以看出，MC 的计算开销与 Q-Learning 相当，但收敛所需的 episode 数更多。

四、实验总结

本次实验实现了值迭代、策略迭代、蒙特卡洛和 Q-Learning 四种强化学习算法，验证了基于模型与无模型方法的核心思想。动态规划方法在已知模型时能够精确求解最优策略，不同奖赏设置显著影响智能体行为；无模型方法中，Q-Learning 凭借单步更新特性展现出更快的收敛速度和更稳定的学习曲线。实验过程中，边界条件处理、状态转移概率计算和探索策略设计是主要实现难点。

A 代码实现

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 import random
5 import time
6
7 # 固定随机种子
8 np.random.seed(42)
9 random.seed(42)
10
11
12 class GridWorld: # 3x4网格世界
13     def __init__(self, step_reward=-0.04):
14         self.rows = 3
15         self.cols = 4
16         self.start = (0, 0)
17         self.wall = (1, 1)
18         self.terminals = {(3, 2): 1.0, (3, 1): -1.0}
19         self.step_reward = step_reward
20         self.gamma = 0.9
21         self.actions = [(0, 1), (0, -1), (-1, 0), (1, 0)]
22         self.action_names = ['^', 'v', '<', '>']
23
24     def is_terminal(self, state):
25         return state in self.terminals
26
27     def get_reward(self, state):
28         if state in self.terminals:
29             return self.terminals[state]
30         return self.step_reward
31
32     def move(self, state, action):
33         new_x = state[0] + action[0]
34         new_y = state[1] + action[1]
35         if new_x < 0 or new_x >= self.cols or \
36             new_y < 0 or new_y >= self.rows or \
37             (new_x, new_y) == self.wall:
```

```
38         return state
39
40
41     def get_transition_probs(self, state, action_idx):
42         probs = []
43         action = self.actions[action_idx]
44         probs.append((0.8, self.move(state, action)))
45         if action[0] == 0:
46             probs.append((0.1, self.move(state, (-1, 0))))
47             probs.append((0.1, self.move(state, (1, 0))))
48         else:
49             probs.append((0.1, self.move(state, (0, 1))))
50             probs.append((0.1, self.move(state, (0, -1))))
51
52         return probs
53
54
55     def all_states(self):
56         states = []
57         for x in range(self.cols):
58             for y in range(self.rows):
59                 if (x, y) != self.wall:
60                     states.append((x, y))
61
62     def value_iteration(env, theta=1e-4): # 值迭代
63         V = {s: 0.0 for s in env.all_states()}
64         delta_history = []
65         time_history = [] # 累计时间
66         start_time = time.time()
67
68         while True:
69             delta = 0
70             for s in env.all_states():
71                 if env.is_terminal(s):
72                     V[s] = env.get_reward(s)
73                     continue
74
75                 v = V[s]
76                 max_val = float('-inf')
77                 for a_idx in range(len(env.actions)):
78                     val = 0
79                     for prob, next_s in env.get_transition_probs(s, a_idx):
80                         val += prob * V[next_s]
81                     max_val = max(max_val, val)
82                 V[s] = env.get_reward(s) + env.gamma * max_val
83                 delta = max(delta, abs(v - V[s]))
84
85             delta_history.append(delta)
86             time_history.append((time.time() - start_time) * 1000) # ms
```

```
86         if delta < theta:
87             break
88
89     # 提取策略
90     policy = {}
91     for s in env.all_states():
92         if env.is_terminal(s):
93             policy[s] = None
94             continue
95         best_a, best_val = -1, float('-inf')
96         for a_idx in range(len(env.actions())):
97             val = sum(p * V[ns] for p, ns in env.get_transition_probs(s, a_idx))
98             if val > best_val:
99                 best_val, best_a = val, a_idx
100             policy[s] = best_a
101
102
103
104 def policy_evaluation(policy, env, V, theta=1e-4): # 策略评估
105     iterations = 0
106     while True:
107         delta = 0
108         for s in env.all_states():
109             if env.is_terminal(s):
110                 continue
111             v = V[s]
112             val = sum(p * V[ns] for p, ns in env.get_transition_probs(s, policy[s]))
113             V[s] = env.get_reward(s) + env.gamma * val
114             delta = max(delta, abs(v - V[s]))
115         iterations += 1
116         if delta < theta:
117             break
118     return V, iterations
119
120
121 def policy_iteration(env): # 策略迭代
122     policy = {s: 0 for s in env.all_states() if not env.is_terminal(s)}
123     for s in env.terminals:
124         policy[s] = None
125     V = {s: 0.0 for s in env.all_states()}
126     for s in env.terminals:
127         V[s] = env.get_reward(s)
128
129     policy_changes = []
130     eval_iterations = []
131     time_history = [] # 累计时间
132     start_time = time.time()
133
```

```

134     while True:
135         V, iters = policy_evaluation(policy, env, V)
136         eval_iterations.append(iters)
137
138         changes = 0
139         for s in env.all_states():
140             if env.is_terminal(s):
141                 continue
142             old_action = policy[s]
143             best_a, best_val = -1, float('-inf')
144             for a_idx in range(len(env.actions)):
145                 val = sum(p * V[ns] for p, ns in env.get_transition_probs(s, a_idx))
146                 if val > best_val:
147                     best_val, best_a = val, a_idx
148             policy[s] = best_a
149             if old_action != best_a:
150                 changes += 1
151
152         policy_changes.append(changes)
153         time_history.append((time.time() - start_time) * 1000)
154         if changes == 0:
155             break
156     return V, policy, policy_changes, eval_iterations, time_history
157
158
159 def monte_carlo_es(env, episodes=10000): # MC探索起点
160     Q = {(s, a): 0.0 for s in env.all_states() for a in range(len(env.actions))}
161     returns = {(s, a): [] for s in env.all_states() for a in range(len(env.actions))}
162     policy = {s: np.random.choice(len(env.actions)) for s in env.all_states() if not env.
163               is_terminal(s)}
164
165     non_terminal = [s for s in env.all_states() if not env.is_terminal(s)]
166     history_returns, q_changes, time_history = [], [], []
167     start_time = time.time()
168
169     for ep in range(episodes):
170         episode = []
171         curr_state = random.choice(non_terminal)
172         curr_action = np.random.choice(len(env.actions))
173
174         steps = 0
175         while not env.is_terminal(curr_state) and steps < 1000:
176             probs = env.get_transition_probs(curr_state, curr_action)
177             r = random.random()
178             cum = 0
179             for p, ns in probs:
180                 cum += p
181                 if r <= cum:

```

```
181         next_state = ns
182         break
183     reward = env.get_reward(next_state)
184     episode.append((curr_state, curr_action, reward))
185     curr_state = next_state
186     if env.is_terminal(curr_state):
187         break
188     curr_action = policy[curr_state]
189     steps += 1
190
191     visited, G, max_change = set(), 0, 0
192     for i in range(len(episode) - 1, -1, -1):
193         s, a, r = episode[i]
194         G = r + env.gamma * G
195         if (s, a) not in visited:
196             visited.add((s, a))
197             old_q = Q[(s, a)]
198             returns[(s, a)].append(G)
199             Q[(s, a)] = np.mean(returns[(s, a)])
200             max_change = max(max_change, abs(Q[(s, a)] - old_q))
201             policy[s] = max(range(len(env.actions)), key=lambda x: Q[(s, x)])
202
203     history_returns.append(sum(r for _, _, r in episode))
204     q_changes.append(max_change)
205     time_history.append((time.time() - start_time) * 1000)
206
207     final_policy = {s: max(range(len(env.actions)), key=lambda a: Q[(s, a)]) for s in
208                     non_terminal}
209
210     return Q, final_policy, history_returns, q_changes, time_history
211
212 def q_learning(env, episodes=10000, alpha=0.1, epsilon=0.1): # Q-Learning
213     Q = {(s, a): 0.0 for s in env.all_states() for a in range(len(env.actions))}

214     non_terminal = [s for s in env.all_states() if not env.is_terminal(s)]
215     history_rewards, q_changes, time_history = [], [], []
216     start_time = time.time()

217     for ep in range(episodes):
218         state = env.start
219         total_reward, max_change = 0, 0
220         steps = 0

221
222         while not env.is_terminal(state) and steps < 1000:
223             if random.random() < epsilon:
224                 action = np.random.choice(len(env.actions))
225             else:
226                 action = max(range(len(env.actions)), key=lambda a: Q[(state, a)])
```

```

228     probs = env.get_transition_probs(state, action)
229     r = random.random()
230     cum = 0
231     for p, ns in probs:
232         cum += p
233         if r <= cum:
234             next_state = ns
235             break
236
237     reward = env.get_reward(next_state)
238     if env.is_terminal(next_state):
239         target = reward
240     else:
241         target = reward + env.gamma * max(Q[(next_state, a)] for a in range(len(env.
242 actions)))
243
243     old_q = Q[(state, action)]
244     Q[(state, action)] += alpha * (target - old_q)
245     max_change = max(max_change, abs(Q[(state, action)] - old_q))
246
247     total_reward += reward
248     state = next_state
249     steps += 1
250
251     history_rewards.append(total_reward)
252     q_changes.append(max_change)
253     time_history.append((time.time() - start_time) * 1000)
254
255 policy = {s: max(range(len(env.actions)), key=lambda a: Q[(s, a)]) for s in non_terminal
256 }
257
258 return Q, policy, history_rewards, q_changes, time_history

259 def plot_grid_policy(policy, V, env, title, filename):
260     plt.figure(figsize=(8, 6))
261     grid_val = np.zeros((3, 4))
262     for s, val in V.items():
263         grid_val[2 - s[1], s[0]] = val
264     for s, reward in env.terminals.items():
265         grid_val[2 - s[1], s[0]] = reward
266     sns.heatmap(grid_val, annot=True, cmap="YlGnBu", cbar=True, fmt='%.2f')
267
268     arrow_map = {0: '↑', 1: '↓', 2: '←', 3: '→'}
269     for s, action in policy.items():
270         if action is not None:
271             plt.text(s[0] + 0.5, 2 - s[1] + 0.25, arrow_map[action],
272                  ha='center', va='center', fontsize=18, color='red', fontweight='bold')
273     plt.text(1.5, 1.5, 'WALL', ha='center', va='center', color='black', fontsize=12,

```

```
backgroundcolor='gray')
274     for s, reward in env.terminals.items():
275         plt.text(s[0] + 0.5, 2 - s[1] + 0.25, '+1' if reward > 0 else '-1',
276                   ha='center', va='center', fontsize=14, color='red', fontweight='bold')
277     plt.title(title)
278     plt.savefig(filename)
279     plt.close()
280
281
282 def plot_rewards(mc_hist, q_hist, filename):
283     plt.figure(figsize=(10, 5))
284     window = 100
285     plt.plot(np.convolve(mc_hist, np.ones(window)/window, mode='valid'), label='Monte Carlo
286 ES')
287     plt.plot(np.convolve(q_hist, np.ones(window)/window, mode='valid'), label='Q-Learning')
288     plt.xlabel('Episodes')
289     plt.ylabel('Episode Return')
290     plt.title('Learning Performance Comparison')
291     plt.legend()
292     plt.grid(True, alpha=0.3)
293     plt.savefig(filename)
294     plt.close()
295
296
297 def plot_convergence(vi_delta, vi_time, pi_changes, pi_evals, pi_time,
298                      mc_changes, mc_time, q_changes, q_time, filename):
299     fig, axes = plt.subplots(2, 2, figsize=(12, 10))
300
301     # 值迭代
302     ax1 = axes[0, 0]
303     ax1.semilogy(vi_delta, 'b-', linewidth=2, label='Delta')
304     ax1.axhline(y=1e-4, color='r', linestyle='--', alpha=0.5)
305     ax1.set_xlabel('Iteration')
306     ax1.set_ylabel('Max Delta (log)', color='blue')
307     ax1.set_title(f'Value Iteration ({len(vi_delta)} iters, {vi_time[-1]:.1f}ms)')
308     ax1.grid(True, alpha=0.3)
309     ax1_t = ax1.twinx()
310     ax1_t.plot(vi_time, 'g--', linewidth=1, label='Time')
311     ax1_t.set_ylabel('Cumulative Time (ms)', color='green')
312
313     # 策略迭代
314     ax2 = axes[0, 1]
315     x = range(1, len(pi_changes) + 1)
316     ax2.bar(x, pi_changes, alpha=0.7, color='blue', label='Policy Changes')
317     ax2.set_xlabel('Round')
318     ax2.set_ylabel('States Changed', color='blue')
319     ax2.set_title(f'Policy Iteration ({len(pi_changes)} rounds, {pi_time[-1]:.1f}ms)')
```

```
320     ax2_t.plot(x, pi_time, 'g-o', linewidth=2, label='Time')
321     ax2_t.set_ylabel('Cumulative Time (ms)', color='green')
322
323     # MC
324     ax3 = axes[1, 0]
325     window = 100
326     mc_smooth = np.convolve(mc_changes, np.ones(window)/window, mode='valid')
327     ax3.semilogy(mc_smooth, 'b-', linewidth=1)
328     ax3.set_xlabel('Episode')
329     ax3.set_ylabel('Max Q Change (log)', color='blue')
330     ax3.set_title(f'Monte Carlo ES ({mc_time[-1]/1000:.2f}s)')
331     ax3.grid(True, alpha=0.3)
332     ax3_t = ax3.twinx()
333     # 每1000个episode画一个时间点
334     time_x = list(range(0, len(mc_time), 1000))
335     time_y = [mc_time[i]/1000 for i in time_x]
336     ax3_t.plot(time_x, time_y, 'g--', linewidth=1)
337     ax3_t.set_ylabel('Time (s)', color='green')
338
339     # Q-Learning
340     ax4 = axes[1, 1]
341     q_smooth = np.convolve(q_changes, np.ones(window)/window, mode='valid')
342     ax4.semilogy(q_smooth, 'b-', linewidth=1)
343     ax4.set_xlabel('Episode')
344     ax4.set_ylabel('Max Q Change (log)', color='blue')
345     ax4.set_title(f'Q-Learning ({q_time[-1]/1000:.2f}s)')
346     ax4.grid(True, alpha=0.3)
347     ax4_t = ax4.twinx()
348     time_x = list(range(0, len(q_time), 1000))
349     time_y = [q_time[i]/1000 for i in time_x]
350     ax4_t.plot(time_x, time_y, 'g--', linewidth=1)
351     ax4_t.set_ylabel('Time (s)', color='green')
352
353     plt.tight_layout()
354     plt.savefig(filename, dpi=150)
355     plt.close()
356
357
358 if __name__ == "__main__":
359     print("Running Value Iteration and Policy Iteration...")
360     env = GridWorld(step_reward=-0.04)
361
362     V_vi, policy_vi, vi_delta, vi_time = value_iteration(env)
363     print(f"  VI: {len(vi_delta)} iterations, {vi_time[-1]:.2f}ms")
364
365     V_pi, policy_pi, pi_changes, pi_evals, pi_time = policy_iteration(env)
366     print(f"  PI: {len(pi_changes)} rounds, {pi_time[-1]:.2f}ms")
```

```
368     # 不同R(s)的策略图
369     for r in [0.01, -0.01, -0.04]:
370         env_r = GridWorld(step_reward=r)
371         V, policy, _, _ = value_iteration(env_r)
372         plot_grid_policy(policy, V, env_r, f"Value Iteration Policy (R={r})", f"policy_vi_{r}.png")
373
374     print("Running MC and Q-Learning...")
375     env = GridWorld(step_reward=-0.04)
376
377     Q_mc, policy_mc, mc_hist, mc_changes, mc_time = monte_carlo_es(env, episodes=10000)
378     print(f"  MC: 10000 episodes, {mc_time[-1]/1000:.2f}s")
379
380     Q_q, policy_q, q_hist, q_changes, q_time = q_learning(env, episodes=10000)
381     print(f"  Q-Learning: 10000 episodes, {q_time[-1]/1000:.2f}s")
382
383     # 画策略图
384     V_mc = {s: max(Q_mc.get((s, a), 0) for a in range(4)) for s in env.all_states() if not env.is_terminal(s)}
385     V_q = {s: max(Q_q.get((s, a), 0) for a in range(4)) for s in env.all_states() if not env.is_terminal(s)}
386
387     plot_grid_policy(policy_mc, V_mc, env, "Monte Carlo Policy", "policy_mc.png")
388     plot_grid_policy(policy_q, V_q, env, "Q-Learning Policy", "policy_q.png")
389     plot_rewards(mc_hist, q_hist, "learning_comparison.png")
390
391     # 收敛过程对比图
392     plot_convergence(vi_delta, vi_time, pi_changes, pi_evals, pi_time,
393                      mc_changes, mc_time, q_changes, q_time, "convergence_comparison.png")
394
395     print("Done.")
```